

- 1 Read data and print few lines
- 2 Basic text processing and cleaning
- 3 Create a corpus
- 4 Transformation of text
- 5 Term Document Matrix & N-Gram Analysis
- 6 Appendix

# Step by Step Text Analytics on My Own Learning Data Log

Code ▾

Ariful Mondal (*ariful dot mondal [at] gmail dot com*)

30 November 2017

This is a quick text analysis using 3 months of my learning completion data on Lynda and LinkedIn Learning using *R*. Some of the required packages are listed below.

Hide

```
setwd("D:\\RProgramming")  
## Load Packages  
library("tm")  
library("stringi")  
library("wordcloud")  
library("clue")  
library("ggplot2")  
library("RColorBrewer")  
library("SnowballC")  
library("RWeka")
```

- **tm**: (<https://cran.r-project.org/web/packages/tm/tm.pdf>) for text mining and natural language processing(NLP)
- **wordcloud**: (<https://cran.r-project.org/web/packages/wordcloud/wordcloud.pdf>) to create word clouds
- **stringi**: (<https://cran.r-project.org/web/packages/stringi/stringi.pdf>) to leverage character string processing facilities
- **clue**: (<https://cran.r-project.org/web/packages/clue/clue.pdf>) for Cluster Ensembles
- **SnowballC**: (<https://cran.r-project.org/web/packages/SnowballC/SnowballC.pdf>) for stemming text
- **RColorBrewer**: (<https://cran.r-project.org/web/packages/RColorBrewer/RColorBrewer.pdf>) color palettes for nice colours of the graphs
- **RWeka**: (<https://cran.r-project.org/web/packages/RWeka/RWeka.pdf>) text mining and N-Gram analysis
- **ggplot2** (<https://cran.r-project.org/package=ggplot2/ggplot2.pdf>) to Create Elegant Data Visualizations Using the Grammar of Graphics

## 1 Read data and print few lines

You may download this data from my google drive here ([https://drive.google.com/open?id=1KudlYfcJ7KbGvrhcSSA4pRpIA3tMswH0w\\_mwxDTs9xQ](https://drive.google.com/open?id=1KudlYfcJ7KbGvrhcSSA4pRpIA3tMswH0w_mwxDTs9xQ))

Hide

```
lnd <- readLines("file:///D:/RProgramming/Ariful_Islam_Mondal_Training_Log_Nov_2017.csv", en  
coding = "UCS-2LE", skipNul = TRUE)  
lnd[1:10]
```

```
## [1] "COURSE, Skills"

## [2] "Learning Information Governance,Document Management"

## [3] "Open Data: Unleashing Hidden Value,\"Big Data, Data Analysis\""

## [4] "Learning Public Data Sets,\"Data Analysis, Microsoft Excel\""

## [5] "Smarter Cities: Using Data to Drive Urban Innovation,\"Big Data, Data Management\""

## [6] "Blockchain Basics,\"Corporate Finance, Databases\""

## [7] "Online Marketing Foundations: Digital Marketing Research,Lead Generation"

## [8] "Marketing Analytics: Presenting Digital Marketing Data,\"Google Analytics, Web Analytics\""

## [9] "Calculating Gross Profit with Google Analytics,\"Google Analytics, E-commerce\""

## [10] "Microsoft Azure for Developers,\"Microsoft Azure, Cloud Development\""
```

## 2 Basic text processing and cleaning

- Remove non-English characters, letters etc. using `iconv()` and option `latin1`
- Remove special characters with spaces using `gsub()` and regular expression `[\^0-9a-z]`
- Remove duplicate characters using `gsub()` and regular expression
- Remove special numbers with spaces using `gsub()` and regular expression
- Remove multiple spaces to one using `gsub()` and regular expression

To know more on `inconv()` click here

(<https://www.rdocumentation.org/packages/base/versions/3.4.1/topics/iconv>), for `gsub()` click here

(<https://www.r-bloggers.com/regular-expression-and-associated-functions-in-r/>) and click here

([https://en.wikipedia.org/wiki/Regular\\_expression](https://en.wikipedia.org/wiki/Regular_expression)) to know more on regular expression, also view `regex`

(<https://stat.ethz.ch/R-manual/R-devel/library/base/html/regex.html>).

Hide

```
# Remove non-English characters, Letters etc.
# Help ?iconv
lnd<-iconv(lnd, "latin1", "ASCII", sub="")
# Remove special characters with spaces
# Help ?gsub
lnd <- gsub("[^\^0-9a-z]", " ", lnd, ignore.case = TRUE)
# Remove duplicate characters
lnd <- gsub('([[:alpha:]]\\1+)', '\\1', lnd)
# Remove special numbers with spaces
lnd <- gsub("[^\^a-z]", " ", lnd, ignore.case = TRUE)
# Remove multiple spaces to one
lnd <- gsub("\\s+", " ", lnd)
lnd <- gsub("^\^\\s", "", lnd)
lnd <- gsub("\\s$", "", lnd)
```

Print after clean up...

Hide

```
# Summary
lnd[1:10]
```

```
## [1] "COURSE Skills"

## [2] "Learning Information Governance Document Management"

## [3] "Open Data Unleashing Hidden Value Big Data Data Analysis"

## [4] "Learning Public Data Sets Data Analysis Microsoft Excel"

## [5] "Smarter Cities Using Data to Drive Urban Innovation Big Data Data Management"

## [6] "Blockchain Basics Corporate Finance Databases"

## [7] "Online Marketing Foundations Digital Marketing Research Lead Generation"

## [8] "Marketing Analytics Presenting Digital Marketing Data Google Analytics Web Analytic
s"
## [9] "Calculating Gross Profit with Google Analytics Google Analytics E commerce"

## [10] "Microsoft Azure for Developers Microsoft Azure Cloud Development"
```

Hide

```
summary(lnd)
```

```
##      Length      Class      Mode
##      115 character character
```

Hide

```
str(lnd)
```

```
## chr [1:115] "COURSE Skills" ...
```

## 3 Create a corpus

Create a virtual corpus using `Vcorpus()` function.

Hide

```
# create Corpus
# Help ??VCorpus
myCorpus <- VCorpus(VectorSource(lnd))
myCorpus
```

```
## <<VCorpus>>
## Metadata: corpus specific: 0, document level (indexed): 0
## Content: documents: 115
```

## 4 Transformation of text

[Optional for already cleaned data]

Perform necessary transformation/preprocessing activities using `tm_map()` from `tm` package. The objective is to have clean texts by removing stop words, punctuation, multiple white spaces etc. We will perform the following transformations

- **tolower:** (<https://www.rdocumentation.org/packages/quanteda/versions/0.99.12/topics/toLower>) normalize text to small cases. For example `My name Is Ariful` will be converted to small case

my name is ariful .

- **stopwords:** (<https://www.rdocumentation.org/packages/tm/versions/0.7-1/topics/stopwords>) remove stop words using English dictionary. For example remove words like “a”, “and”, “but”, “how”, “or”, and “what”
- **removePunctuation:** (<https://www.rdocumentation.org/packages/tm/versions/0.7-1/topics/removePunctuation>) remove punctuation marks such as ! " # \$ % & ' ( ) \* + , - . / : ; < = > ? @ [ ] ^ \_ ` { | } ~.
- **stripWhitespace:** (<https://www.rdocumentation.org/packages/tm/versions/0.7-1/topics/stripWhitespace>) Strip extra white space from a text document. Multiple white space characters are collapsed to a single blank
- **PlainTextDocument:** (<https://www.rdocumentation.org/packages/tm/versions/0.7-1/topics/PlainTextDocument>) convert document to plain text format

[Hide](#)

```
# Help ??tm_map'

# Normalize to small cases
myCorpus <- tm_map(myCorpus, content_transformer(tolower))

# Remove Stop Words
myCorpus <- tm_map(myCorpus, removeWords, stopwords("english"))

# Remove Punctuation
myCorpus <- tm_map(myCorpus, removePunctuation)

# Remove Numbers
myCorpus <- tm_map(myCorpus, removeNumbers)

# Create plain text documents
myCorpus <- tm_map(myCorpus, PlainTextDocument)

# Stem words in a text document using Porter's stemming algorithm.
myCorpus <- tm_map(myCorpus, stemDocument, "english")

# Strip White Spaces
myCorpus <- tm_map(myCorpus, stripWhitespace)
```

## 5 Term Document Matrix & N-Gram Analysis

Now we will use `TermDocumentMatrix()` to create a document-term matrix or term-document matrix which is a mathematical matrix that describes the frequency of terms/words/strings that occur in a collection of documents. In a document-term matrix, rows correspond to documents in the collection and columns correspond to terms. There are various schemes for determining the value that each entry in the matrix should take. Read more on wiki ([https://en.wikipedia.org/wiki/Document-term\\_matrix](https://en.wikipedia.org/wiki/Document-term_matrix)).

In the fields of computational linguistics and probability, an n-gram is a contiguous sequence of n items from a given sequence of text or speech. The items can be phonemes, syllables, letters, words or base pairs according to the application.

The n-grams typically are collected from a text or speech corpus. When the items are words, n-grams may also be called shingles.

- An n-gram of size 1 is referred to as a "unigram" ;
- An n-gram of size 2 is a "bigram" (or, less commonly, a "digram");
- An n-gram of size 3 is a "trigram" .

Larger sizes are sometimes referred to by the value of n in modern language, e.g., “four-gram”, “five-gram”, and so on. [wiki (<https://en.wikipedia.org/wiki/N-gram>)].

### 5.1 Create Unigram

[Hide](#)


```
unitdm <- TermDocumentMatrix(myCorpus)
mat <- as.matrix(unitdm)
wf <- sort(rowSums(mat),decreasing=TRUE)
df <- data.frame(word = names(wf),freq=wf)
head(df, 10)
```

```
##           word freq
## data      data  67
## management management 36
## learning   learning 33
## azure      azure  27
## microsoft  microsoft 27
## leadership leadership 26
## business   business  23
## foundations foundations 23
## cloud      cloud  22
## analysis   analysis  19
```

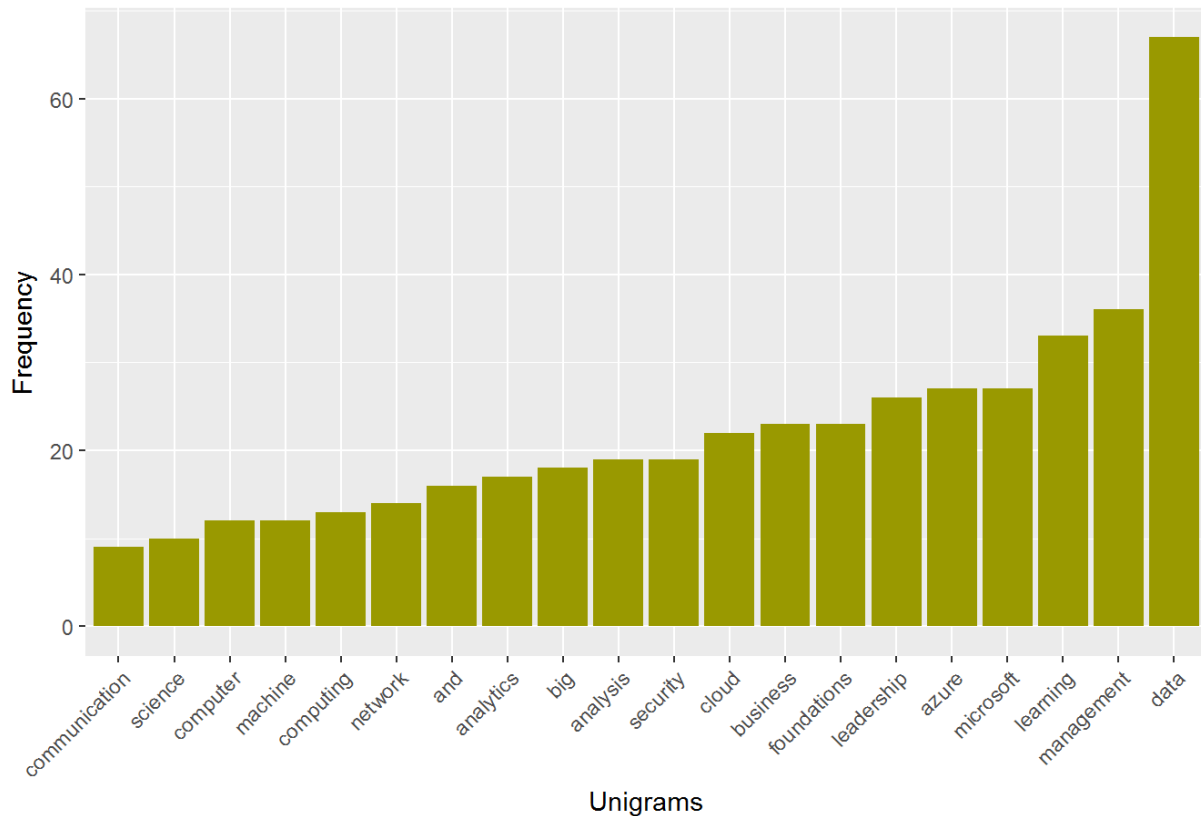
## 5.1.1 Ploting unigram

[Hide](#)

```
barplot(df[1:20,]$freq, las = 2, names.arg = df[1:20,]$word,
        col =df[1:20,]$freq, main = "",
        ylab = "Frequencies", cex.axis=.8, cex = .8, cex.lab=0.75, cex.main=.75)
```

[Hide](#)


```
ggplot(df[1:20,], aes(x = reorder(df[1:20,]$word, df[1:20,]$freq), y = df[1:20,]$freq)) +
  geom_bar(stat = "identity", fill = "#999900") +
  labs(title = " ") +
  xlab("Unigrams") +
  ylab("Frequency")+
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



## 5.1.2 Create Word cloud with unigram

[Hide](#)

```
set.seed(1234)
wordcloud(words = df$word, freq = df$freq, min.freq = 1,
  max.words=100, random.order=FALSE, rot.per=0.75, las=3,
  colors=brewer.pal(8, "Dark2"), c(5,.5), vfont=c("script","plain"))
```



### 5.1.3 Find frequently occurred words

[Hide](#)

```
findFreqTerms(unitdm, lowfreq = 5)
```

```
## [1] "administration" "analysis"      "analytics"     "and"
## [5] "azure"          "big"           "business"      "cloud"
## [9] "communication"  "computer"      "computing"     "data"
## [13] "databases"      "decision"      "design"         "development"
## [17] "essential"      "excel"         "executive"     "finance"
## [21] "for"            "foundations"   "google"        "hadoop"
## [25] "intelligence"   "leadership"    "learning"      "machine"
## [29] "management"     "marketing"     "microsoft"     "modeling"
## [33] "network"        "operations"    "science"       "security"
## [37] "skills"         "statistics"    "training"      "web"
## [41] "with"
```

### 5.1.4 Find Associations between Words

Find associations in document-term or term-document matrix using function `findAssocs(x, terms, corlimit)` from *tm* package.

- **x**: A DocumentTermMatrix or a TermDocumentMatrix.
- **terms**: a character vector holding terms.
- **corlimit**: a numeric vector (of the same length as terms; recycled otherwise) for the (inclusive) lower correlation limits of each term in the range from zero to one.

1. Find associated words with *data*...

[Hide](#)

```
findAssocs(unitdm, terms = "data", corlimit = 0.35)
```

```
## $data
##      big      science      modeling      analysis      career
##      0.71      0.56      0.52      0.48      0.48
## certifications database      paths      steps      intelligence
##      0.48      0.48      0.48      0.48      0.40
##      analytics visualization
##      0.38      0.35
```

2. Find associated words with *machine*...

Hide

```
findAssocs(unitdm, terms = "machine", corlimit = 0.35)
```

```
## $machine
##      learning      trees estimations mathematica      python
##      0.68      0.49      0.47      0.47      0.38
```

3. Find associated words with *management*...

Hide

```
findAssocs(unitdm, terms = "management", corlimit = 0.35)
```

```
## $management
## operations      small      nonprofit      finding      high potentials
##      0.62      0.51      0.41      0.35      0.35      0.35
##      records      retaining
##      0.35      0.35
```

4. Find associated words with *azure*...

Hide

```
findAssocs(unitdm, terms = "azure", corlimit = 0.35)
```

```
## $azure
##      microsoft      computing      implement      cloud      networking
##      0.85      0.58      0.53      0.52      0.52
## administration      active      directory      virtual      development
##      0.51      0.40      0.40      0.36      0.35
```

5. Find associated words with *security*...

Hide

```
findAssocs(unitdm, terms = "security", corlimit = 0.35)
```

```
## $security
##      computer      network      cert      compliance      comptia
##      0.78      0.68      0.53      0.53      0.53
##      operational      prep      asset      cissp      cryptography
##      0.53      0.53      0.40      0.40      0.40
## investigation      response
##      0.40      0.40
```

## 5.2 Crate Bigram

Hide



```
BigramTokenizer <- function(x) NGramTokenizer(x, Weka_control(min = 2, max = 2)) # Create bigram tokenizer using RWeka
bitdm <- TermDocumentMatrix(myCorpus, control = list(tokenize = BigramTokenizer)) # Create bigram
inspect(bitdm[15:30,1:20]) # Inspect few bigrams
```

```
## <<TermDocumentMatrix (terms: 16, documents: 20)>>
## Non-/sparse entries: 6/314
## Sparsity : 98%
## Maximal term length: 20
## Weighting : term frequency (tf)
## Sample :
## Docs
## Terms      1 14 19 2 3 4 5 6 7 8
## ai advanced    0  0  0 0 0 0 0 0 0 0
## ai foundations 0  0  0 0 0 0 0 0 0 0
## amazon web      0  0  0 0 0 0 0 0 0 0
## analysis business 0  0  0 0 0 0 0 0 0 0
## analysis data    0  1  0 0 0 0 0 0 0 0
## analysis microsoft 0  0  0 0 0 1 0 0 0 0
## analysis office  0  0  1 0 0 0 0 0 0 0
## analysis web      0  1  0 0 0 0 0 0 0 0
## analytics business 0  1  0 0 0 0 0 0 0 0
## analytics career 0  1  0 0 0 0 0 0 0 0
```

Hide


```
mat_bigram <- as.matrix(bitdm)
wf_bigram <- sort(rowSums(mat_bigram),decreasing=TRUE)
df1 <- data.frame(word = names(wf_bigram),freq=wf_bigram)
head(df1, 10)
```

```
## word freq
## microsoft azure      microsoft azure 21
## big data              big data 18
## data analysis         data analysis 14
## cloud computing       cloud computing 13
## machine learning      machine learning 12
## data science          data science 10
## computer security     computer security 8
## computing microsoft   computing microsoft 8
## network security      network security 8
## essential training    essential training 7
```

## 5.2.1 Ploting Bigram


Hide

```
biplot<-barplot(df1[1:20,]$freq, las = 2, names.arg = df1[1:20,]$word,
  col = df1[1:20,]$freq, main="",
  ylab = "Frequencies", cex.axis=.65, cex = .65, cex.lab=0.5, cex.main=.75)
```



Hide

```
ggplot(df1[1:20,], aes(x = reorder(df1[1:20,]$word, df1[1:20,]$freq), y = df1[1:20,]$freq)) +
  geom_bar(stat = "identity", fill = "#00b3b3") +
  labs(title = " ") +
  xlab("Bigrams") +
  ylab("Frequency")+
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



Hide



Hide

```
## <<TermDocumentMatrix (terms: 16, documents: 20)>>
## Non-/sparse entries: 6/314
## Sparsity      : 98%
## Maximal term length: 29
## Weighting      : term frequency (tf)
## Sample        :
##
##              Docs
## Terms
## ai advanced decision      0 0 0 0 0 0 0 0 0 0
## ai foundations decision   0 0 0 0 0 0 0 0 0 0
## ai foundations value      0 0 0 0 0 0 0 0 0 0
## amazon web services       0 0 0 0 0 0 0 0 0 0
## analysis data management  0 1 0 0 0 0 0 0 0 0
## analysis microsoft excel  0 0 0 0 0 1 0 0 0 0
## analysis office web       0 0 1 0 0 0 0 0 0 0
## analysis web analytics    0 1 0 0 0 0 0 0 0 0
## analytics business analysis 0 1 0 0 0 0 0 0 0 0
## analytics career paths    0 1 0 0 0 0 0 0 0 0
```

Hide


```
mat_trigram <- as.matrix(tritdm)
wf_trigram <- sort(rowSums(mat_trigram),decreasing=TRUE)
df2 <- data.frame(word = names(wf_trigram),freq=wf_trigram)
head(df2, 10)
```

```
##
##              word freq
## cloud computing microsoft      8
## computing microsoft azure      8
## computer security network      7
## security network security      7
## administration computer networking 4
## azure network administration    4
## big data data                  4
## microsoft azure network         4
## network administration computer 4
## amazon web services             3
```

### 5.3.1 Ploting Trigram


Hide

```
triplot<-barplot(df2[1:20,]$freq, las = 2, names.arg = df2[1:20,]$word,
  col = df1[1:20,]$freq, main="",
  ylab = "Frequencies", cex.axis=.65, cex = .65, cex.lab=0.5, cex.main=.75)
```



Hide

```
ggplot(df2[1:20,], aes(x = reorder(df2[1:20,]$word, df2[1:20,]$freq), y = df2[1:20,]$freq)) +  
  geom_bar(stat = "identity", fill=df2[1:20,]$freq) +  
  labs(title = " ") +  
  xlab("Trigrams") +  
  ylab("Frequency")+  
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



## 5.3.2 Create Word cloud with Trigram

Hide

```
set.seed(1234)
wordcloud(words = df2$word, freq = df2$freq, min.freq = 3,
          max.words=100, random.order=T, rot.per=0.75,
          colors=brewer.pal(8, "Dark2"), c(1.7,.7), vfont=c("script","plain"))
```



More coming soon....

## 6 Appendix

### 6.1 About R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com> (<http://rmarkdown.rstudio.com>).

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

#### 6.1.1 Including summary

Hide

```
summary(cars)
```